

Multi-Card Retrieval: Purpose-Specific Embedding Views for Aspect-Rich Corpora

Anonymous

Anonymous

Abstract. A document that covers several distinct matters is usually indexed as one dense vector, which averages those matters together. We quantify the resulting loss and show how to avoid it cheaply. On a controlled corpus in which each document carries exactly k near-orthogonal aspects, the similarity between an aspect-directed query and a pooled document vector decays as $k^{-0.32}$ while the maximum over per-aspect embeddings stays flat, and retrieval quality follows: at $k = 10$, nDCG@10 falls from 0.682 to 0.273 under pooling but holds at 0.672 under per-aspect cards. On the LIMIT capacity benchmark, where single-vector retrieval is known to fail, the decisive comparison is against chunking rather than against pooling: splitting the same text into fixed windows lifts nDCG@10 by 0.076, while splitting it into purpose-aligned cards lifts it by 0.674. The benefit therefore comes from aligning the unit of representation with the unit the query asks about, not from storing more vectors per document. We report the storage multiplier each corpus actually incurs, and we release the code, the seeds, and every per-query score.

Keywords: Dense retrieval · Multi-vector representations · Aspect-based retrieval · Reproducibility

1 Introduction

Enterprise corpora are full of documents that are about several things at once.

Contributions.

1. A measurement of aspect dilution on a controlled corpus, giving a fitted decay exponent that is milder than the orthogonality argument predicts (Section 3).
2. Evidence that purpose alignment, not embedding count, drives the benefit, via a chunking control on identical source text (Section 4).
3. An open, seeded, single-command reproduction of every number reported.

2 Related work

3 Aspect dilution

We first measure the effect in isolation, on a corpus built so that the number of aspects per document is a controlled variable.

Table 1. Aspect dilution. Deltas are per-query paired differences, card system minus pooled, with 95% bootstrap intervals and paired permutation p -values Holm-corrected across all 21 comparisons. Relevance-set sizes grow with k by construction, so nDCG is comparable down a column and between the two systems within a row, but not across rows.

	k	pooled	cards	Δ	95% CI	win/tie/loss
1	0.682	0.682	+0.000		exact tie	0/200/0
2	0.395	0.571	+0.177	[+0.145, +0.207]		160/0/40
3	0.308	0.565	+0.257	[+0.226, +0.288]		173/1/26
4	0.271	0.556	+0.285	[+0.253, +0.317]		175/0/25
5	0.300	0.585	+0.285	[+0.251, +0.318]		176/2/22
7	0.282	0.656	+0.374	[+0.339, +0.408]		182/0/18
10	0.273	0.672	+0.399	[+0.360, +0.438]		186/1/13

Construction. Each document concatenates k passages drawn from k of ten mutually unrelated topic pools. Every passage realises one *fact*, a (pool, attribute, action) triple, and a query asks about one fact in wording that shares no template with the passages. Relevance is therefore defined by the generative process rather than assumed: every document carrying a passage that realises the queried fact is relevant, which yields relevance sets of 3.7 documents per query at $k = 1$ rising to 33.3 at $k = 10$. Two representations of the same document are compared: one pooled embedding of the whole text, and k per-aspect embeddings scored by their maximum. Both are scored at document level. We use 500 documents and 200 queries per setting, encoder `all-MiniLM-L6-v2`, seed 13.

Result. Table 1 and Figure 1 report the outcome. The per-aspect representation holds its similarity to relevant documents almost exactly constant as k grows, from 0.498 at $k = 1$ to 0.493 at $k = 10$, while the pooled representation decays monotonically to 0.235. Retrieval follows: the pooled system loses more than half its nDCG@10 between $k = 1$ and $k = 2$ and does not recover, while the card system stays between 0.556 and 0.682 throughout. At $k = 1$ the two systems are identical by construction, and the measurement confirms it exactly, which is a useful check on the pipeline.

The theory has the right shape and the wrong magnitude. Treating a pooled vector as the normalised mean of k near-orthogonal aspect vectors predicts that an aspect-directed query’s similarity falls as $1/\sqrt{k}$. The measured ratio of pooled to card similarity instead follows $k^{-0.32}$, well above the prediction at every k . Two properties of real encoders explain the gap, and both work in the same direction: aspects drawn from natural language are not orthogonal, since all of the passages share register and syntax, and encoder similarities are anisotropic, occupying a narrow cone that compresses differences. We therefore report the orthogonality argument as an upper bound on dilution rather than as a prediction that the data confirm.




Fig. 1. Aspect dilution as a function of k . Left: the aspect signal is preserved by cards and lost by pooling. Centre: measured dilution against the $1/\sqrt{k}$ prediction and a fitted power law. Right: the retrieval consequence.

Table 2. LIMIT-small. Deltas are against the pooled dense baseline, paired permutation with Holm correction over 1,000 queries. Chunking and carding both raise the embedding count from one per document to many; only carding recovers the lost performance.

system	nDCG@10	R@2	R@10	Δ vs pooled p
BM25	0.997	0.991	1.000	+0.683 < 0.001
dense, pooled	0.314	0.162	0.480	baseline
dense, 30-word chunks	0.390	0.225	0.549	+0.076 < 0.001
dense, per-aspect cards	0.988	0.965	0.997	+0.674 < 0.001
cards $\oplus$ BM25 (RRF)	0.997	0.990	1.000	+0.683 < 0.001

4 Capacity: does alignment or count explain the gain?

The previous section compares cards against a single pooled vector, which leaves an obvious objection open: cards store more vectors per document, so perhaps capacity alone explains the gain and any decomposition would do. This section answers that objection with a control.

Setting. LIMIT [1] is constructed to expose the representational ceiling of single-vector retrieval. Each document names a person and lists the things they like; each query asks who likes one specific thing. The small variant has 46 documents, 44.0 attributes each, and 1,000 queries, so the task looks trivial and yet defeats dense retrievers. It also hands us an unambiguous aspect decomposition, which makes it the sharpest available test of whether alignment matters. We compare five systems, all scored at document level: BM25; a pooled dense vector; the same text cut into 30-word windows with 10-word overlap (186 chunks); the same text cut into per-attribute cards (2,024 cards); and reciprocal rank fusion of the card system with BM25.

Alignment, not count. Table 2 separates the two explanations. Chunking multiplies the embedding budget and recovers 0.076 nDCG@10. Carding multiplies

the same budget and recovers 0.674, roughly nine times as much, from identical source text under an identical encoder. The difference between the two is only where the boundaries fall: a 30-word window spans several attributes and dilutes each of them in the same way the pooled vector does, whereas a card corresponds to exactly the unit a query names. Capacity is necessary and not sufficient; the decomposition has to match the query distribution.

What this result does not show. LIMIT supplies a perfect decomposition for free, and its queries are effectively exact-attribute lookups, which is why BM25 wins outright and the hybrid only matches it. Real corpora arrive undecomposed, the taxonomy must be designed and may be misaligned, and the honest reading of this table is as an upper bound on what alignment can buy rather than as an estimate of what it will buy. We also note the cost: 44 cards per document here, far above the small constant that the pattern implies elsewhere, so the storage multiplier is a property of the corpus and must be reported per corpus rather than assumed.

5 The first pass: does a cheap gate earn its place?

A two-pass design rests on a first pass that is cheap enough to run over everything and discriminating enough to be worth running. This section tests that premise, and finds that it holds while two pieces of received practice around it do not.

Setting. On-purpose items are 4,000 messages from a public corporate email corpus; off-purpose items are 4,000 Usenet posts from unrelated groups. Provenance gives exact labels at no annotation cost. Three anchor sets describe facets of the purpose, each item scores the mean of its best few similarities to each set, and the combination is fitted on one half and reported on the other.

The gate discriminates. Held-out ROC-AUC is 0.933 and PR-AUC 0.935. A first pass built from nothing but embeddings and a handful of anchor phrases separates on-purpose from off-purpose material well enough to be a filter rather than a sample.

Fitting the weights is worth almost nothing. Giving the three anchor sets equal weight yields ROC-AUC 0.931 against 0.933 for fitted weights. Practitioners tune such weightings at length; on this evidence the discrimination comes from having several purpose-specific anchor sets at all, not from how they are mixed.

The filter is not where the saving comes from. Table 3 corrects a common framing. A first pass is usually described as discarding most of a corpus, but at recall 0.95 it discards about a third, and at recall 0.99 about an eighth. Gating therefore contributes a linear factor between one and two, bought directly with recall. The order-of-magnitude saving in a two-pass design comes from elsewhere: from attaching generative cost to the number of discovered topics rather than to the number of documents. Attributing the saving to the filter overstates the filter and understates the substitution.

Table 3. What the gate discards, by recall target. The economic case for a first pass is proportional to the fraction it removes, and at the high recall such a filter is supposed to hold, that fraction is modest.

recall target	achieved recall	precision	corpus discarded
0.80	0.781	0.910	57.7%
0.90	0.901	0.830	46.5%
0.95	0.954	0.710	33.7%
0.98	0.981	0.605	20.0%
0.99	0.993	0.557	12.1%

Table 4. Selection policies on a constructed coverage task. All coverage gains over relevance ranking are significant under Holm-corrected paired permutation tests. Relevance differences are within noise, an artefact of the strong gate discussed below.

policy	S-recall	α -nDCG	nDCG@10	answerless rate
relevance only	0.628	0.692	0.965	0.033
MMR, $\lambda = 0.7$	0.666	0.723	0.964	0.032
MMR, $\lambda = 0.3$	0.818	0.840	0.963	0.036
DPP	0.879	0.877	0.963	0.035
cluster round robin	0.666	0.724	0.967	0.028
outlier harvest, 20%	0.716	0.734	0.966	0.031
outlier harvest, 40%	0.772	0.775	0.967	0.028

Caveat. Separating corporate email from Usenet is easier than separating on-purpose from off-purpose material inside one corpus, which is what a deployed gate faces. These numbers bound discrimination from above; failure here would have been disqualifying, and success here is necessary rather than sufficient.

6 Selection: how much diversity, and for whom

The preceding sections concern how a document is represented. This one concerns what is selected once candidates are retrieved, which is where the study’s second claim lives: that the right amount of diversity depends on who consumes the results. This section settles the measurable half of that claim, namely what each policy actually delivers; the consumer half is future work (Section 7).

Setting. Each of 120 constructed queries is associated with several distinct facts, and each candidate document realises at most one of them, so the subtopics a selection covers can be counted exactly rather than inferred. Deriving subtopics by clustering would have been circular, since one of the policies under test selects by cluster. Candidates are gated to the upper half by relevance first, as a first pass would do, and each policy then selects ten items.

The operator we proposed is not the best operator. Table 4 is not the result we expected. Harvesting the low-similarity tail of a gated pool does raise coverage substantially, by 0.144 S-recall at a 40 percent quota, but a determinantal point process raises it by 0.251 and maximal marginal relevance at $\lambda = 0.3$ by 0.190. We report this because it sharpens rather than weakens the position we are arguing for. The claim worth defending is not that a particular diversification operator is best; it is that the quantity of diversity is a function of the consumer, and that function applies to whichever operator a system already has. On this evidence the strongest implementation of the idea is a determinantal point process whose quality-diversity balance is set per consumer, not the harvest rule.

Under a strong gate, coverage is nearly free. No policy paid a measurable relevance penalty, and the answerless rate did not move. Both follow from gating: when the pool is already bounded for relevance, spending part of the budget away from its centre costs little and risks little. This is an argument for the gate rather than evidence that diversity is free, and we would expect the trade to reappear with a weaker filter or a stricter relevance criterion.

7 Limitations

Three limitations bound what has been shown so far.

First, both results here are on constructed corpora. The synthetic corpus fixes k by construction, and LIMIT is adversarial by design. Neither tells us what happens on prose written by people for their own purposes, which is the setting the pattern is meant for.

Second, the aspect decomposition is given rather than inferred in both cases. The practical question is whether a taxonomy designed in advance, and applied by an imperfect extractor, retains a useful share of the benefit.

Third, a single encoder is used throughout. Anisotropy differs between encoders, and the dilution exponent is likely to move with it.

8 Reproducibility

Every number in this paper is produced by one command per experiment against committed code and seeds, and every per-query score is released alongside the aggregates so that the significance tests can be recomputed independently.

Determinism is enforced rather than assumed: two seeded runs must produce byte-identical per-query output, which is checked in the repository. Reaching that required accumulating similarity in double precision, because on some platforms a single-precision matrix product varies by around 10^{-8} between runs with the reduction order, enough to move a recorded value across a rounding boundary.

Metric implementations are covered by hand-worked golden tests, and the selection policies by property tests, for example that maximal marginal relevance with $\lambda = 1$ reduces exactly to relevance ranking and that reciprocal rank fusion is invariant to the order of its input lists.

References

1. Weller, O., Boratko, M., Naim, I., Lee, J.: On the theoretical limitations of embedding-based retrieval. arXiv preprint arXiv:2508.21038 (2025)